

R Exercise

This is a three-part exercise. In all three, the goal is to estimate the parameters of a model. None of this will be graded, but you are encouraged to challenge yourself. If you do not have much programming experience, the first part may be a bit difficult; if that is the case, don't get disappointed and invest your time on Parts 2 and 3.

In all three parts, you can use the following lines (and change them to your liking) to create your own data and then see how well your functions are performing compared to the outputs you get from R's built in functions (lm and glm).

```
library(MASS)
N = 100
#creating independent variables
### first, the covariance matrix for our multi-variate normal
### this determines the correlation of our independent variables.
sig = matrix(c(2,.5,.25,.5,1,0,.25,0,1) , nrow=3)
### now the variables:
M = mvrnorm(n = N, mu = rep(1,3), Sigma = sig )

y.cont = 1 + 2* M[,1] - 5 * M[,2] + M[,3] + rnorm(N)
y.bin = as.numeric ( y.cont > 0 )

### include the intercept in your independent variables
X = cbind (1, M )
y = y.cont ### or y = y.bin , depending on the exercise
```

Part 1: OLS using a recursive function (optional)

We want to write a function (call it *reg*) that takes two inputs (y: dependent variable, X: a matrix whose columns are independent variables) and returns a list which contains its two outputs (b: estimated betas, and e: residuals).

$$y = b_1 x[,1] + b_2 x[,2] + \dots + e$$

The task is to write this function in a recursive way (an example is provided below). Every time you call the function, it should break the X matrix into two parts and call itself for each part in appropriate ways and then stitch the results back together in a correct way.

The only time when the parameters are directly calculated is when x has only one column. In this case, you can calculate the parameter by simply doing

$$b = \text{sum}(y*x) / \text{sum}(x^2)$$

In other situations, you simply use the following rule.

b[n] is equal to the result of regressing the **residual from regression of y on X[, -n]** on the **residual from regression of X[,n] on X[, -n]**

Example of a recursive function:

The following function calculates x^i for integer values of i

```
pwr = function(x,i){
## if i==0, return 1
## if i>0, return x^(i-1) * x
## if i<0, return x^(i+1) / x
if (as.integer(i) == 0 ) return( 1)
else return(
      ifelse( i>0, pwr(x,i-1)* x , pwr(x,i+1) /x ))
}
# test
pwr(10 , -3) ### should be 0.001
pwr(2 , 5)   ### should be 32
```

Part 2: OLS using numerical optimization

We want to write a function (call it *regress*) that takes two inputs (y: dependent variable, X: a matrix whose columns are independent variables) and returns estimated coefficients.

$$y = b_1 x[,1] + b_2 x[,2] + \dots + e$$

The task is to write this function using `optim`. Follow these steps:

- Write a function, call it *RSS*, that takes *b*, *X*, and *y* (*b* should be first) and returns the total sum of squared residuals.
- Define *b0* and make it such that it could be a good starting point.
- Use `optim` within *regress* to find optimal values of *b*. Notice that “...” in the R’s usage definition means “every other argument you want to pass on to the function” so you can call `optim` as you wish and then just type *X=X* and *y=y* and it would pass these to your *RSS* function. Here’s `optim`’s usage definition.

```
optim(par, fn, gr = NULL, ...,
      method = c("Nelder-Mead", "BFGS", "CG", "L-BFGS-B", "SANN", "Brent"),
      lower = -Inf, upper = Inf, control = list(), hessian = FALSE)
```

For a given *b*, you can calculate *y.hat* using matrix multiplication (instead of a loop):

```
y.hat = X %*% b
```

but if you are using this, you have to make sure that *X* is a matrix with correct dimensions.

Part 3: Logit

This is similar to Part 2, but we are going to maximize the likelihood function for the logistic regression. You should write a function named *logit* that takes *y* (binary dependent variable) and *X* (a matrix whose columns are the independent variables).

$$Pr(y = 1) = \frac{1}{1 + \exp(-X \times b)}$$

Before we do that, let us go through a simple explanation and example of maximum likelihood estimation (MLE).

Suppose we have a model about how the world works but we do not know the true value of the parameters (*b*) of the model. There is also some source of randomness in the world.

We make observations from the world.

We calculate how likely each observation would be in the world if the parameters were *b*.

We take the logarithm of the likelihoods. (only for mathematical convenience).

We add all of these log-likelihood's for the observations.

The logic of MLE is that the best estimate for the unknown parameters are values of *b* that maximize the likelihood of observing what we have observed.

This is quite simple, but it takes a while to sink in. The following example should demonstrate the logic.

Example-----

Assume that we have a coin. The coin has an unknown parameter called tailiness which is the likelihood of observing tails when we toss it. Let *b* represent tailiness of our coin.

We make *N* observations of coin tosses. Each observation is coded as "1" if it is tails and "0" if it is heads. Let *y_i* represent this.

Likelihood of observing 1= *b*

Likelihood of observing 0= 1-*b*

`Ly = ifelse(y==1, b, 1-b)`

Log Likelihood of observing the entire vector of *y* (assuming that coin tosses are independent of each other)

`LLy = sum (log (Ly))`

Now we can find the optimal solution:

```

# create the data first for tailiness=.333:
y = rbinom(n=100,p=.333,size=1)

LLy = function(b, y) return(
    sum(
        log (ifelse(y==1, b, 1-b) ) ))

b0 = 0.5 #starting value
### we are doing maximization with optim this time
optim(b0, fn = LLy, y=y, control = list( fnscale = -1 ) )

###ignore the warning that tells you the default method is
### not the best method for the problem at hand.

### If there is convergence, optim's output $par should be
### close to 0.333
-----

```

Now, we want to use MLE to estimate a logistic regression.

Assume that the parameter to be estimated is b , independent variables form matrix X and observed binary outcomes are y . Using the following steps, write a function $LL(b,y,X)$ which is then going to be maximized.

- What is the likelihood that y_i is 1?
- What is the likelihood that y_i is 0?
- What is the sum of log-likelihood's of observing all of y ?
- Pick a good starting point for b
- Use optim to find the value of b that maximizes the likelihood of observing what we have observed.
- report this value as b if there is convergence