

MT0.

Which of the following statements about map-reduce are true? Check all that apply.

- (a) If you only have 1 computer with 1 computing core, then map-reduce is unlikely to help
- (b) If we run map-reduce using N computers, then we will always get at least an N-Fold speedup
- (c) Because of network latency and other overhead associated with map-reduce, if we run r N-Fold speedup compared to using 1 computer
- (d) When using map-reduce with gradient descent, we usually use a single machine that accumulates updates from all machines, in order to compute the parameter update for the iteration

A, C, D

MT1.

Suppose you wish to write a MapReduce job that creates normalized word co-occurrence counts. If many reducers receive appropriate normalization factors (denominators) in the correct order (ignoring overhead), the mapper should emit according to which pattern:

- (a) emit (*,word) count
- (b) There is no need to use order inversion here
- (c) emit (word,*) count
- (d) None of the above

C

MT2.

When searching for frequent itemsets with the Apriori algorithm (using a threshold, N), the A occurrences of the itemset $\{A,B,C\}$ provided

- (a) all subsets of $\{A,B,C\}$ occur less than N times.
- (b) any pair of $\{A,B,C\}$ occurs less than N times.
- (c) any subset of $\{A,B,C\}$ occurs less than N times.
- (d) All of the above

C

MT3.

When building a Bayesian document classifier, Laplace smoothing serves what purpose?

- (a) It allows you to use your training data as your validation data.
- (b) It prevents zero-products in the posterior distribution.
- (c) It accounts for words that were missed by regular expressions.
- (d) None of the above

B

MT4.

By increasing the complexity of a model regressed on some samples of data, it is likely that

- (a) Increased variance and bias
- (b) Increased variance and decreased bias
- (c) Decreased variance and bias
- (d) Decreased variance and increased bias

B

MT5.

Combiners can be integral to the successful utilization of the Hadoop shuffle. This utility is a

- (a) minimization of reducer workload
- (b) both (a) and (c)
- (c) minimization of network traffic
- (d) none of the above

B

===Pairwise similarity using K-L divergence===

In probability theory and information theory, the Kullback–Leibler divergence (also informally KL divergence) is a non-symmetric measure of the difference between two probability distributions. The KL divergence of Q from P, denoted $DKL(P||Q)$, is a measure of the information lost when Q is used to approximate P.

For discrete probability distributions P and Q, the Kullback–Leibler divergence of Q from P is

$$KLDistance(P, Q) = \sum_i (P(i) \log (P(i) / Q(i)))$$

In the extreme cases, the KL Divergence is 1 when P and Q are maximally different and is 0 when P and Q are the same distribution).

For more information on K-L Divergence see:

https://en.wikipedia.org/wiki/Kullback%E2%80%93Leibler_divergence (https://en.wikipedia.org/wiki/Kullback%E2%80%93Leibler_divergence)

For the next three questions we will use an MRJob class for calculating pairwise similarity using K-L Divergence.

Job 1: create inverted index (assume just two objects)

Job 2: calculate/accumulate the similarity of each pair of objects using K-L Divergence

Download the following notebook and then fill in the code for the first reducer to calculate the similarity of line1 and line2, i.e., $KLD(line1||line2)$.

Here we ignore characters which are not alphabetical. And all alphabetical characters are lower case.

<http://nbviewer.ipython.org/urls/dl.dropbox.com/s/9onx4c2dujtkgd7/Kullback%E2%80%93Leibler%20divergence%20notebook.ipynb>
<http://nbviewer.ipython.org/urls/dl.dropbox.com/s/9onx4c2dujtkgd7/Kullback%E2%80%93Leibler%20divergence%20notebook.ipynb>
<https://www.dropbox.com/s/zr9xfhwakrxz9hc/Kullback%E2%80%93Leibler%20divergence%20notebook.ipynb>
<https://www.dropbox.com/s/zr9xfhwakrxz9hc/Kullback%E2%80%93Leibler%20divergence%20notebook.ipynb>

```
In [1]: %%writefile kltext.txt
1.Data Science is an interdisciplinary field about processes and systems t
2.Machine learning is a subfield of computer science[1] that evolved from

Writing kltext.txt
```

```
In [2]: import numpy as np
np.log(3)
```

```
Out[2]: 1.0986122886681098
```

```

In [1]: %%writefile kldivergence.py
from mrjob.job import MRJob
import re
import numpy as np
class kldivergence(MRJob):
    def mapper1(self, _, line):
        index = int(line.split('.',1)[0])
        letter_list = re.sub(r"^[A-Za-z]+", '', line).lower()
        count = {}
        for l in letter_list:
            if count.has_key(l):
                count[l] += 1
            else:
                count[l] = 1
        for key in count:
            yield key, [index, (count[key]+1)*1.0/(len(letter_list)+24)]

    def reducer1(self, key, values):
        objects = {}
        if key in ['p','t']: print 'not A'
        elif key in ['k','q']: print 'not B'
        elif key in ['j','q']: print 'not C'
        elif key in ['j','f']: print 'not C'
        for i, p in values:
            objects[i] = p
        yield None, objects[1]*np.log(objects[1] / objects[2])

    def reducer2(self, key, values):
        kl_sum = 0
        for value in values:
            kl_sum = kl_sum + value
        yield None, kl_sum

    def steps(self):
        return [self.mr(mapper=self.mapper1,
                        reducer=self.reducer1),
                self.mr(reducer=self.reducer2)]

if __name__ == '__main__':
    kldivergence.run()

```

Overwriting kldivergence.py

```
In [2]: import kldivergence
from kldivergence import kldivergence
mr_job = kldivergence(args=['kltext.txt'])
with mr_job.make_runner() as runner:
    runner.run()
    # stream_output: get access of the output
    for line in runner.stream_output():
        print mr_job.parse_output_line(line)
```

WARNING:mrjob.runner:

WARNING:mrjob.runner:PLEASE NOTE: Starting in mrjob v0.5.0, protocols will n your job with --strict-protocols or set up mrjob.conf as described at ht ady-for-strict-protocols (<https://pythonhosted.org/mrjob/whats-new.html#re>

WARNING:mrjob.runner:

WARNING:mrjob.job:mr() is deprecated and will be removed in v0.6.0. Use mr

WARNING:mrjob.job:mr() is deprecated and will be removed in v0.6.0. Use mr

WARNING:mrjob.job:mr() is deprecated and will be removed in v0.6.0. Use mr

WARNING:mrjob.job:mr() is deprecated and will be removed in v0.6.0. Use mr

WARNING:mrjob.job:mr() is deprecated and will be removed in v0.6.0. Use mr

WARNING:mrjob.job:mr() is deprecated and will be removed in v0.6.0. Use mr

WARNING:mrjob.job:mr() is deprecated and will be removed in v0.6.0. Use mr

WARNING:mrjob.job:mr() is deprecated and will be removed in v0.6.0. Use mr

WARNING:mrjob.job:mr() is deprecated and will be removed in v0.6.0. Use mr

WARNING:mrjob.job:mr() is deprecated and will be removed in v0.6.0. Use mr

WARNING:mrjob.job:mr() is deprecated and will be removed in v0.6.0. Use mr

WARNING:mrjob.job:mr() is deprecated and will be removed in v0.6.0. Use mr

WARNING:mrjob.job:mr() is deprecated and will be removed in v0.6.0. Use mr

WARNING:mrjob.job:mr() is deprecated and will be removed in v0.6.0. Use mr

WARNING:mrjob.job:mr() is deprecated and will be removed in v0.6.0. Use mr

WARNING:mrjob.job:mr() is deprecated and will be removed in v0.6.0. Use mr

WARNING:mrjob.job:mr() is deprecated and will be removed in v0.6.0. Use mr

WARNING:mrjob.job:mr() is deprecated and will be removed in v0.6.0. Use mr

not C

not C

not B

not A

not B

not A

(None, 0.06823525136041805)

MT6.

Which number below is the closest to the result you get for KLD(Line1||line2)?

(a) 0.7

(b) 0.5

(c) 0.2

(d) 0.1

D

MT7.

Which of the following letters are missing from these character vectors?

(a) p and t

(b) k and q

(c) j and q

(d) j and f

D

MT8. The KL divergence on multinomials is defined only when they have nonzero entr

For zero entries, we have to smooth distributions. Suppose we smooth in this way:

$$(n_i + 1)/(n + 24)$$

where n_i is the count for letter i and n is the total count of all letters. After smoothing, which $KLD(\text{Line1} || \text{line2})$??

(a) 0.08

(b) 0.71

(c) 0.02

(d) 0.11

A

MT9.

Which of the following are true statements with respect to gradient descent for machine learning?

- (a) To make gradient descent converge, we must slowly decrease alpha over time and use a
- (b) Gradient descent is guaranteed to find the global minimum for any function $J()$ regardless
- (c) Gradient descent can converge even if alpha is kept fixed. (But alpha cannot be too large speed up the process.
- (d) For the specific choice of cost function $J()$ used in linear regression, there is no local opti

C, D

Write a MapReduce job in MRJob to do the training at scale of a weighted K-means algorithm

You can write your own code or you can use most of the code from the following notebook:

<http://nbviewer.ipython.org/urls/dl.dropbox.com/s/kjtdyi10nwmk4ko/MrJobKmeans-MIDS->
<http://nbviewer.ipython.org/urls/dl.dropbox.com/s/kjtdyi10nwmk4ko/MrJobKmeans-MIDS->
<https://www.dropbox.com/s/kjtdyi10nwmk4ko/MrJobKmeans-MIDS-Midterm.ipynb?dl=0>
<https://www.dropbox.com/s/kjtdyi10nwmk4ko/MrJobKmeans-MIDS-Midterm.ipynb?dl=0>

Weight each example as follows using the inverse vector length (Euclidean norm):

$\text{weight}(X) = 1/\|X\|$,

where $\|X\| = \text{SQRT}(X.X) = \text{SQRT}(X_1^2 + X_2^2)$

Here X is vector made up of X1 and X2.

Using the following data answer the following questions:

<https://www.dropbox.com/s/ai1uc3q2ucverly/Kmeandata.csv?dl=0> (<https://www.dropbox.c>

```
In [3]: %%writefile Kmeans.py
from numpy import argmin, array, random, sqrt
from mrjob.job import MRJob
from mrjob.step import MRJobStep
from itertools import chain
DIR = '/Users/bshur/School/ML at Scale/MT/'

#Calculate find the nearest centroid for data point
def MinDist(datapoint, centroid_points):
    datapoint = array(datapoint)
    centroid_points = array(centroid_points)
    diff = datapoint - centroid_points
    diffsq = diff**2
```



```

distances = (diffsq.sum(axis = 1))*0.5
# Get the nearest centroid for each instance
min_idx = argmin(distances)
return min_idx

#Check whether centroids converge
def stop_criterion(centroid_points_old, centroid_points_new,T):
    oldvalue = list(chain(*centroid_points_old))
    newvalue = list(chain(*centroid_points_new))
    Diff = [abs(x-y) for x, y in zip(oldvalue, newvalue)]
    Flag = True
    for i in Diff:
        if(i>T):
            Flag = False
            break
    return Flag

class MRKmeans(MRJob):
    centroid_points=[]
    k=3
    def steps(self):
        return [
            MRJobStep(mapper_init = self.mapper_init
                        , mapper=self.mapper
                        ,combiner = self.combiner
                        ,reducer=self.reducer)
        ]
    #load centroids info from file
    def mapper_init(self):
        self.centroid_points = [map(float,s.split('\n')[0].split(',')) for
                                open(DIR+'Centroids.txt', 'w').close()]
    #load data and output the nearest centroid index and data point
    def mapper(self, _, line):
        D = (map(float,line.split(',')))
        idx = MinDist(D,self.centroid_points)
        weight = 1/sqrt(D[0]**2 + D[1]**2)
        yield int(idx), (D[0],D[1],1*weight)
    #Combine sum of data points locally
    def combiner(self, idx, inputdata):
        sumx = sumy = num = 0
        for x,y,n in inputdata:
            num = num + n
            sumx = sumx + x
            sumy = sumy + y
        yield int(idx), (sumx,sumy,num)
    #Aggregate sum for each cluster and then calculate the new centroids
    def reducer(self, idx, inputdata):
        centroids = []
        num = [0]*self.k
        distances = 0
        for i in range(self.k):

```

```
for i in range(self.k):
    centroids.append([0,0])
for x, y, n in inputdata:
    num[idx] = num[idx] + n
    centroids[idx][0] = centroids[idx][0] + x
    centroids[idx][1] = centroids[idx][1] + y
centroids[idx][0] = centroids[idx][0]/num[idx]
centroids[idx][1] = centroids[idx][1]/num[idx]
with open(DIR+'Centroids.txt', 'a') as f:
    f.writelines(str(centroids[idx][0]) + ',' + str(centroids[idx]
yield idx,(centroids[idx][0],centroids[idx][1]))

if __name__ == '__main__':
    MRKmeans.run()
```

Overwriting Kmeans.py

```

In [5]: from numpy import random, array
        from Kmeans import MRKmeans, stop_criterion
        mr_job = MRKmeans(args=['Kmeandata.csv'])

        #Generate initial centroids
        centroid_points = [[5.777968353788965672e+00,1.179139375692149772e-01]
                           ,[8.451051977473833077e+00,-2.377148039960867987e-01]
                           ,[3.903195518555621080e-01,5.495947017581701566e+00]]

        k = 3
        with open('/Users/bshur/School/ML at Scale/MT/Centroids.txt', 'w+') as f:
            f.writelines(','.join(str(j) for j in i) + '\n' for i in centroid_

        # Update centroids iteratively
        for i in range(10):
            # save previous centroids to check convergency
            centroid_points_old = centroid_points[:]
            print "iteration"+str(i+1)+": "
            with mr_job.make_runner() as runner:
                runner.run()
                # stream_output: get access of the output
                for line in runner.stream_output():
                    key,value = mr_job.parse_output_line(line)
                    print key, value
                    centroid_points[key] = value

            print "\n"
            i = i + 1
        print "Centroids\n"
        print centroid_points

WARNING:mrjob.runner:
WARNING:mrjob.step:MRJobStep has been renamed to MRStep. The old name will
WARNING:mrjob.step:MRJobStep has been renamed to MRStep. The old name will
WARNING:mrjob.step:MRJobStep has been renamed to MRStep. The old name will
WARNING:mrjob.step:MRJobStep has been renamed to MRStep. The old name will
WARNING:mrjob.step:MRJobStep has been renamed to MRStep. The old name will
WARNING:mrjob.step:MRJobStep has been renamed to MRStep. The old name will
WARNING:mrjob.step:MRJobStep has been renamed to MRStep. The old name will

[0.14865107226547203, 7.98472075745883]

iteration10:
0 [0.14865107226547203, 7.98472075745883]

Centroids

[[0.14865107226547203, 7.98472075745883], [0.14865107226547203, 7.98472075
09]]

```

MT10.

Which result below is the closest to the centroids you got after running your weighted K-me

- (a) (-4.0,0.0), (4.0,0.0), (6.0,6.0)
- (b) (-4.5,0.0), (4.5,0.0), (0.0,4.5)
- (c) (-5.5,0.0), (0.0,0.0), (3.0,3.0)
- (d) (-4.5,0.0), (-4.0,0.0), (0.0,4.5)

C

MT11.

Using the result of the previous question, which number below is the closest to the average assigned (closest) centroid? The average weighted distance is defined as sum over i (weigh

- (a) 2.5
- (b) 1.5
- (c) 0.5
- (d) 4.0

B

MT12.

Which of the following statements are true? Select all that apply.

- a) Since K-Means is an unsupervised learning algorithm, it cannot overfit the data, and thus as is computationally feasible.
- b) The standard way of initializing K-means is setting $\mu_1 = \dots = \mu_k$ to be equal to a vector of zeros.
- c) For some datasets, the "right" or "correct" value of K (the number of clusters) can be aml carefully at the data to decide.
- d) A good way to initialize K-means is to select K (distinct) examples from the training set as examples.

C, D

